

Python for Data Analysis

Time Series

Time Series Data

- *Timestamps*, specific instants in time
- Fixed *periods*, such as the month January 2007 or the full year 2010
- *Intervals* of time, indicated by a start and end timestamp. Periods can be thought of as special cases of intervals
- Experiment or elapsed time; each timestamp is a measure of time relative to a particular start time. For example, the diameter of a cookie baking each second since being placed in the oven

Basic Settings

```
1 import numpy as np
2 import pandas as pd
3 np.random.seed(12345)
4 import matplotlib.pyplot as plt
5 plt.rc('figure', figsize=(10, 6))
6 PREVIOUS_MAX_ROWS = pd.options.display.max_rows
7 pd.options.display.max_rows = 20
8 np.set_printoptions(precision=4, suppress=True)
```

Date and Time Data Types and Tools

```
1 from datetime import datetime
2 now = datetime.now()
3 now
4 now.year, now.month, now.day
```

(2019, 3, 1)

Date and Time Data Types and Tools

```
1 delta = datetime(2011, 1, 7) - datetime(2008, 6, 24, 8, 15)
2 delta
3 delta.days
4 delta.seconds
```

56700

Date and Time Data Types and Tools

```
1 from datetime import timedelta
2 start = datetime(2011, 1, 7)
3 start + timedelta(12)
4 start - 2 * timedelta(12)
```

```
datetime.datetime(2010, 12, 14, 0, 0)
```

Date and Time Data Types and Tools

Types in datetime module

Type	Description
<code>date</code>	Store calendar date (year, month, day) using the Gregorian calendar.
<code>time</code>	Store time of day as hours, minutes, seconds, and microseconds
<code>datetime</code>	Stores both date and time
<code>timedelta</code>	Represents the difference between two datetime values (as days, seconds, and microseconds)

Converting Between String and Datetime

```
1 stamp = datetime(2011, 1, 3)
2 str(stamp)
3 stamp.strftime('%Y-%m-%d')
```

'2011-01-03'

Converting Between String and Datetime

Datetime format specification (ISO C89 compatible)

Type	Description
%Y	4-digit year
%y	2-digit year
%m	2-digit month [01, 12]
%d	2-digit day [01, 31]
%H	Hour (24-hour clock) [00, 23]
%I	Hour (12-hour clock) [01, 12]
%M	2-digit minute [00, 59]
%S	Second [00, 61] (seconds 60, 61 account for leap seconds)
%w	Weekday as integer [0 (Sunday), 6]

Converting Between String and Datetime

Datetime format specification (ISO C89 compatible)

Type	Description
%U	Week number of the year [00, 53]. Sunday is considered the first day of the week, and days before the first Sunday of the year are "week 0".
%W	Week number of the year [00, 53]. Monday is considered the first day of the week, and days before the first Monday of the year are "week 0".
%z	UTC time zone offset as +HHMM or -HHMM, empty if time zone naive
%F	Shortcut for %Y-%m-%d, for example 2012-4-18
%D	Shortcut for %m/%d/%y, for example 04/18/12

Converting Between String and Datetime

```
1 value = '2011-01-03'
2 datetime.strptime(value, '%Y-%m-%d')
3 datestrs = ['2011/7/6', '2011/8/6']
4 [datetime.strptime(x, '%Y/%m/%d') for x in datestrs]
```

```
[datetime.datetime(2011, 7, 6, 0, 0), datetime.datetime(2011, 8, 6,
```

Converting Between String and Datetime

```
1 from dateutil.parser import parse
2 parse('2011-01-03')
```

```
datetime.datetime(2011, 1, 3, 0, 0)
```

```
1 parse('Jan 31, 1997 10:45 PM')
```

```
datetime.datetime(1997, 1, 31, 22, 45)
```

Converting Between String and Datetime

```
1 datestrs = ['2011-07-06 12:00:00', '2011-08-06 00:00:00']  
2 pd.to_datetime(datestrs)
```

```
DatetimeIndex(['2011-07-06 12:00:00', '2011-08-06 00:00:00'], dtype='datetime64[ns]',  
freq=None)
```

```
1 idx = pd.to_datetime(datestrs + [None])  
2 idx
```

```
DatetimeIndex(['2011-07-06', '2011-08-06', 'NaT'], dtype='datetime64[ns]', freq=None)
```

Converting Between String and Datetime

```
1 idx[2]
```

NaT

```
1 pd.isnull(idx)
```

array([False, False, True])

Converting Between String and Datetime

Locale-specific date formatting

Type	Description
%a	Abbreviated weekday name
%A	Full weekday name
%b	Abbreviated month name
%B	Full month name
%c	Full date and time, for example 'Tue 01 May 2012 04:20:57 PM'
%p	Locale equivalent of AM or PM
%x	Locale-appropriate formatted date; e.g. in US May 1, 2012 yields '05/01/2012'
%X	Locale-appropriate time, e.g. '04:24:12 PM'

Time Series Basics

```
1 from datetime import datetime
2 dates = [datetime(2011, 1, 2), datetime(2011, 1, 5),
3          datetime(2011, 1, 7), datetime(2011, 1, 8),
4          datetime(2011, 1, 10), datetime(2011, 1, 12)]
5 ts = pd.Series(np.random.randn(6), index=dates)
6 ts
```

```
2011-01-02    -1.078419
2011-01-05     2.213262
2011-01-07     0.447194
2011-01-08    -0.099447
2011-01-10    -0.573992
2011-01-12     1.727834
dtype: float64
```

Time Series Basics

```
1 ts.index
```

```
DatetimeIndex(['2011-01-02', '2011-01-05', '2011-01-07', '2011-01-08',  
              '2011-01-10', '2011-01-12'],  
              dtype='datetime64[ns]', freq=None)
```

```
1 ts + ts[::-2]
```

```
2011-01-02    -2.156838  
2011-01-05         NaN  
2011-01-07     0.894388  
2011-01-08         NaN  
2011-01-10    -1.147984  
2011-01-12         NaN  
dtype: float64
```

Time Series Basics

```
1 ts.index.dtype
```

```
dtype('<M8[ns]')
```

```
1 stamp = ts.index[0]
```

```
2 stamp
```

```
Timestamp('2011-01-02 00:00:00')
```

Indexing, Selection, Subsetting

```
1 stamp = ts.index[2]  
2 ts[stamp]
```

0.4471938539950634

```
1 ts['2011/1/10']  
2 ts['20110110']
```

-0.5739919980605656

Indexing, Selection, Subsetting

```
1 longer_ts = pd.Series(np.random.randn(1000),  
2                       index=pd.date_range('2000/1/1', periods=1000))  
3 longer_ts  
4 longer_ts['2001']
```

2001-01-01 -0.151545

2001-01-02 0.401587

2001-01-03 -2.223506

2001-01-04 -0.574654

2001-01-05 0.786210

...

2001-12-29 0.405850

2001-12-30 0.680251

2001-12-31 1.357221

Freq: D, Length: 365, dtype: float64

Indexing, Selection, Subsetting

```
1 longer_ts['2001-05']
```

```
2001-05-01    0.734799
```

```
2001-05-02   -0.163056
```

```
2001-05-03    0.196350
```

```
2001-05-04    0.727743
```

```
2001-05-05    0.161154
```

```
...
```

```
2001-05-28   -0.548465
```

```
2001-05-29    0.909617
```

```
2001-05-30    1.475824
```

```
2001-05-31    0.584234
```

```
Freq: D, dtype: float64
```

Indexing, Selection, Subsetting

```
1 ts[datetime(2011, 1, 7):]
```

```
2011-01-07    -0.519439
```

```
2011-01-08    -0.555730
```

```
2011-01-10     1.965781
```

```
2011-01-12     1.393406
```

```
dtype: float64
```

Indexing, Selection, Subsetting

1	ts
---	----

```
2011-01-02    -1.078419
2011-01-05     2.213262
2011-01-07     0.447194
2011-01-08    -0.099447
2011-01-10    -0.573992
2011-01-12     1.727834
dtype: float64
```

1	ts['2011/1/6':'2011/1/11']
---	----------------------------

```
2011-01-07     0.447194
2011-01-08    -0.099447
2011-01-10    -0.573992
dtype: float64
```


Indexing, Selection, Subsetting

```
1 ts.truncate(after='2011/1/9')
```

```
2011-01-02    -1.078419
```

```
2011-01-05     2.213262
```

```
2011-01-07     0.447194
```

```
2011-01-08    -0.099447
```

```
dtype: float64
```

Indexing, Selection, Subsetting

```
1 dates = pd.date_range('2000/1/1', periods=100, freq='W-WED')
2 long_df = pd.DataFrame(np.random.randn(100, 4),
3                         index=dates,
4                         columns=['Colorado', 'Texas',
5                                'New York', 'Ohio'])
6 long_df.loc['2001-5']
```

	Colorado	Texas	New York	Ohio
2001-05-02	-0.661257	1.134968	0.060154	-0.630441
2001-05-09	-2.303203	1.135953	0.039481	1.492634
2001-05-16	0.219628	-2.433629	-0.071871	-1.438643
2001-05-23	-0.963054	0.689989	-1.173536	0.536754
2001-05-30	0.745970	-0.169002	-0.585304	-0.793187

Time Series with Duplicate Indices

```
1 dates = pd.DatetimeIndex(['2000/1/1', '2000/1/2', '2000/1/2',  
2                             '2000/1/2', '2000/1/3'])  
3 dup_ts = pd.Series(np.arange(5), index=dates)  
4 dup_ts
```

```
2000-01-01    0  
2000-01-02    1  
2000-01-02    2  
2000-01-02    3  
2000-01-03    4  
dtype: int32
```

Time Series with Duplicate Indices

```
1 dup_ts.index.is_unique
```

False

```
1 dup_ts['2000/1/3'] # not duplicated
```

4

```
1 dup_ts['2000/1/2'] # duplicated
```

```
2000-01-02    1
2000-01-02    2
2000-01-02    3
dtype: int32
```

Time Series with Duplicate Indices

```
1 grouped = dup_ts.groupby(level=0)
2 grouped.mean()
```

```
2000-01-01    0
2000-01-02    2
2000-01-03    4
dtype: int32
```

```
1 grouped.count()
```

```
2000-01-01    1
2000-01-02    3
2000-01-03    1
dtype: int64
```

Date Ranges, Frequencies, and Shifting

```
1 ts
```

```
2011-01-02    -1.078419
2011-01-05     2.213262
2011-01-07     0.447194
2011-01-08    -0.099447
2011-01-10    -0.573992
2011-01-12     1.727834
dtype: float64
```

```
1 resampler = ts.resample('D')
2 resampler
```

```
DatetimeIndexResampler [freq=<Day>, axis=0, closed=left, label=left, convention=start, base=0]
```

Generating Date Ranges

```
1 index = pd.date_range('2012-04-01', '2012-06-01')
2 index
```

```
DatetimeIndex(['2012-04-01', '2012-04-02', '2012-04-03', '2012-04-04',
               '2012-04-05', '2012-04-06', '2012-04-07', '2012-04-08',
               '2012-04-09', '2012-04-10', '2012-04-11', '2012-04-12',
               '2012-04-13', '2012-04-14', '2012-04-15', '2012-04-16',
               ...,
               '2012-05-19', '2012-05-20', '2012-05-21', '2012-05-22',
               '2012-05-23', '2012-05-24', '2012-05-25', '2012-05-26',
               '2012-05-27', '2012-05-28', '2012-05-29', '2012-05-30',
               '2012-05-31', '2012-06-01'],
              dtype='datetime64[ns]', freq='D')
```

Generating Date Ranges

```
1 pd.date_range(start='2012-04-01', periods=20)
```

```
DatetimeIndex(['2012-04-01', '2012-04-02', '2012-04-03', '2012-04-04',  
               '2012-04-05', '2012-04-06', '2012-04-07', '2012-04-08',  
               '2012-04-09', '2012-04-10', '2012-04-11', '2012-04-12',  
               '2012-04-13', '2012-04-14', '2012-04-15', '2012-04-16',  
               '2012-04-17', '2012-04-18', '2012-04-19', '2012-04-20'],  
              dtype='datetime64[ns]', freq='D')
```


Generating Date Ranges

```
1 pd.date_range(end='2012-06-01', periods=20)
```

```
DatetimeIndex(['2012-05-13', '2012-05-14', '2012-05-15', '2012-05-16',  
               '2012-05-17', '2012-05-18', '2012-05-19', '2012-05-20',  
               '2012-05-21', '2012-05-22', '2012-05-23', '2012-05-24',  
               '2012-05-25', '2012-05-26', '2012-05-27', '2012-05-28',  
               '2012-05-29', '2012-05-30', '2012-05-31', '2012-06-01'],  
              dtype='datetime64[ns]', freq='D')
```

Generating Date Ranges

```
1 pd.date_range('2000-01-01', '2000-12-01', freq='BM')
```

```
DatetimeIndex(['2000-01-31', '2000-02-29', '2000-03-31', '2000-04-28',  
               '2000-05-31', '2000-06-30', '2000-07-31', '2000-08-31',  
               '2000-09-29', '2000-10-31', '2000-11-30'],  
              dtype='datetime64[ns]', freq='BM')
```

Generating Date Ranges

Base Time Series Frequencies

Alias	Offset Type	Description
D	Day	Calendar daily
B	BusinessDay	Business daily
H	Hour	Hourly
T or min	Minute	Minutely
S	Second	Secondly
L or ms	Milli	Millisecond (1/1000th of 1 second)
U	Micro	Microsecond (1/1000000th of 1 second)
M	MonthEnd	Last calendar day of month
BM	BusinessMonthEnd	Last business day (weekday) of month
MS	MonthBegin	First calendar day of month

Generating Date Ranges

Base Time Series Frequencies

Alias	Offset Type	Description
BMS	BusinessMonthBegin	First weekday of month
W-MON, W-TUE, ...	Week	Weekly on given day of week: MON, TUE, WED, THU, FRI, SAT, or SUN.
WOM-1MON, WOM-2MON, ...	WeekOfMonth	Generate weekly dates in the first, second, third, or fourth week of the month. For example, WOM-3FRI for the 3rd Friday of each month.
Q-JAN, Q-FEB, ...	QuarterEnd	Quarterly dates anchored on last calendar day of each month, for year ending in indicated month: JAN, FEB, MAR, APR, MAY, JUN, JUL, AUG, SEP, OCT, NOV, or DEC.
BQ-JAN, BQ-FEB, ...	BusinessQuarterEnd	Quarterly dates anchored on last weekday day of each month, for year ending in indicated month
QS-JAN, QS-FEB, ...	QuarterBegin	Quarterly dates anchored on first calendar day of each month, for year ending in indicated month

Generating Date Ranges

Base Time Series Frequencies

Alias	Offset Type	Description
BQS-JAN, BQS-FEB, ...	BusinessQuarterBegin	Quarterly dates anchored on first weekday day of each month, for year ending in indicated month
A-JAN, A-FEB, ...	YearEnd	Annual dates anchored on last calendar day of given month: JAN, FEB, MAR, APR, MAY, JUN, JUL, AUG, SEP, OCT, NOV, or DEC.
BA-JAN, BA-FEB, ...	BusinessYearEnd	Annual dates anchored on last weekday of given month
AS-JAN, AS-FEB, ...	YearBegin	Annual dates anchored on first day of given month
BAS-JAN, BAS-FEB, ...	BusinessYearBegin	Annual dates anchored on first weekday of given month

Generating Date Ranges

```
1 pd.date_range('2012-05-02 12:56:31', periods=5)
```

```
DatetimeIndex(['2012-05-02 12:56:31', '2012-05-03 12:56:31',  
              '2012-05-04 12:56:31', '2012-05-05 12:56:31',  
              '2012-05-06 12:56:31'],  
              dtype='datetime64[ns]', freq='D')
```

```
1 pd.date_range('2012-05-02 12:56:31', periods=5, normalize=True)
```

```
DatetimeIndex(['2012-05-02', '2012-05-03', '2012-05-04', '2012-05-05',  
              '2012-05-06'],  
              dtype='datetime64[ns]', freq='D')
```

Frequencies and Date Offsets

```
1 from pandas.tseries.offsets import Hour, Minute
2 hour = Hour()
3 hour
```

<Hour>

```
1 four_hours = Hour(4)
2 four_hours
```

<4 * Hours>

Frequencies and Date Offsets

```
1 pd.date_range('2000-01-01', '2000-01-03 23:59', freq='4h')
```

```
DatetimeIndex(['2000-01-01 00:00:00', '2000-01-01 04:00:00',  
               '2000-01-01 08:00:00', '2000-01-01 12:00:00',  
               '2000-01-01 16:00:00', '2000-01-01 20:00:00',  
               '2000-01-02 00:00:00', '2000-01-02 04:00:00',  
               '2000-01-02 08:00:00', '2000-01-02 12:00:00',  
               '2000-01-02 16:00:00', '2000-01-02 20:00:00',  
               '2000-01-03 00:00:00', '2000-01-03 04:00:00',  
               '2000-01-03 08:00:00', '2000-01-03 12:00:00',  
               '2000-01-03 16:00:00', '2000-01-03 20:00:00'],  
              dtype='datetime64[ns]', freq='4H')
```


Frequencies and Date Offsets

```
1 Hour(2) + Minute(30)
```

<150 * Minutes>

```
1 pd.date_range('2000-01-01', periods=10, freq='1h30min')
```

```
DatetimeIndex(['2000-01-01 00:00:00', '2000-01-01 01:30:00',  
               '2000-01-01 03:00:00', '2000-01-01 04:30:00',  
               '2000-01-01 06:00:00', '2000-01-01 07:30:00',  
               '2000-01-01 09:00:00', '2000-01-01 10:30:00',  
               '2000-01-01 12:00:00', '2000-01-01 13:30:00'],  
              dtype='datetime64[ns]', freq='90T')
```

Week of month dates

```
1 rng = pd.date_range('2012-01-01', '2012-09-01', freq='WOM-3FRI')
2 list(rng)
```

```
[Timestamp('2012-01-20 00:00:00', freq='WOM-3FRI'),
 Timestamp('2012-02-17 00:00:00', freq='WOM-3FRI'),
 Timestamp('2012-03-16 00:00:00', freq='WOM-3FRI'),
 Timestamp('2012-04-20 00:00:00', freq='WOM-3FRI'),
 Timestamp('2012-05-18 00:00:00', freq='WOM-3FRI'),
 Timestamp('2012-06-15 00:00:00', freq='WOM-3FRI'),
 Timestamp('2012-07-20 00:00:00', freq='WOM-3FRI'),
 Timestamp('2012-08-17 00:00:00', freq='WOM-3FRI')]
```

Shifting (Leading and Lagging) Data

```
1 ts = pd.Series(np.random.randn(4),  
2                 index=pd.date_range('2000/1/1', periods=4, freq='M'))  
3 ts
```

```
2000-01-31    -0.150786  
2000-02-29     1.502825  
2000-03-31    -0.539127  
2000-04-30     1.044848  
Freq: M, dtype: float64
```

Shifting (Leading and Lagging) Data

```
1 ts.shift(2)
```

```
2000-01-31      NaN
2000-02-29      NaN
2000-03-31    -0.150786
2000-04-30     1.502825
Freq: M, dtype: float64
```

```
1 ts.shift(-2)
```

```
2000-01-31    -0.539127
2000-02-29     1.044848
2000-03-31      NaN
2000-04-30      NaN
Freq: M, dtype: float64
```

computing percent changes in a time series or multiple timeseries as DataFrame columns.

```
ts / ts.shift(1) - 1
```

Shifting (Leading and Lagging) Data

```
1 ts.shift(2, freq='M')
```

```
2000-03-31    -0.150786
2000-04-30     1.502825
2000-05-31    -0.539127
2000-06-30     1.044848
Freq: M, dtype: float64
```

```
1 ts.shift(3, freq='D')
```

```
2000-02-03    -0.150786
2000-03-03     1.502825
2000-04-03    -0.539127
2000-05-03     1.044848
dtype: float64
```

```
1 ts.shift(1, freq='90T')
```

```
2000-01-31 01:30:00    -0.150786
2000-02-29 01:30:00     1.502825
2000-03-31 01:30:00    -0.539127
2000-04-30 01:30:00     1.044848
Freq: M, dtype: float64
```

Shifting dates with offsets

```
1 from pandas.tseries.offsets import Day, MonthEnd
2 now = datetime(2011, 11, 17)
3 now + 3 * Day()
```

Timestamp('2011-11-20 00:00:00')

```
1 now + MonthEnd()
```

Timestamp('2011-11-30 00:00:00')

```
1 now + MonthEnd(2)
```

Timestamp('2011-12-31 00:00:00')

Shifting dates with offsets

```
1 offset = MonthEnd()  
2 offset.rollforward(now)
```

```
Timestamp('2011-11-30 00:00:00')
```

```
1 offset.rollback(now)
```

```
Timestamp('2011-10-31 00:00:00')
```

Shifting dates with offsets

```
1 ts = pd.Series(np.random.randn(20),  
2                 index=pd.date_range('2000/1/15', periods=20, freq='4d'))  
3 ts  
4 ts.groupby(offset.rollforward).mean()
```

```
2000-01-31    -0.370374  
2000-02-29     0.335717  
2000-03-31    -0.677309  
dtype: float64
```


Shifting dates with offsets

```
1 ts.resample('M').mean()
```

```
2000-01-31    -0.370374
```

```
2000-02-29     0.335717
```

```
2000-03-31    -0.677309
```

```
Freq: M, dtype: float64
```

Periods and Period Arithmetic

```
p = pd.Period(2007, freq='A-DEC')  
p
```

```
Period('2007', 'A-DEC')
```

```
p + 5  
p - 2
```

```
Period('2005', 'A-DEC')
```

```
pd.Period('2014', freq='A-DEC') - p
```

```
<7 * YearEnds: month=12>
```

Periods and Period Arithmetic

```
rng = pd.period_range('2000-01-01', '2000-06-30', freq='M')  
rng
```

```
PeriodIndex(['2000-01', '2000-02', '2000-03', '2000-04', '2000-05', '2000-06'], dtype='period[M]', freq='M')
```

```
pd.Series(np.random.randn(6), index=rng)
```

```
2000-01    2.089154  
2000-02   -0.060220  
2000-03   -0.167933  
2000-04    0.631634  
2000-05   -1.594313  
2000-06   -1.519937  
Freq: M, dtype: float64
```

Periods and Period Arithmetic

```
values = ['2001Q3', '2002Q2', '2003Q1']  
index = pd.PeriodIndex(values, freq='Q-DEC')  
index
```

```
PeriodIndex(['2001Q3', '2002Q2', '2003Q1'], dtype='period[Q-DEC]', freq='Q-DEC')
```

Period Frequency Conversion

```
1 p = pd.Period('2007', freq='A-DEC')  
2 p  
3 p.asfreq('M', how='start')
```

Period('2007-01', 'M')

```
1 p.asfreq('M', how='end')
```

Period('2007-12', 'M')

Period Frequency Conversion

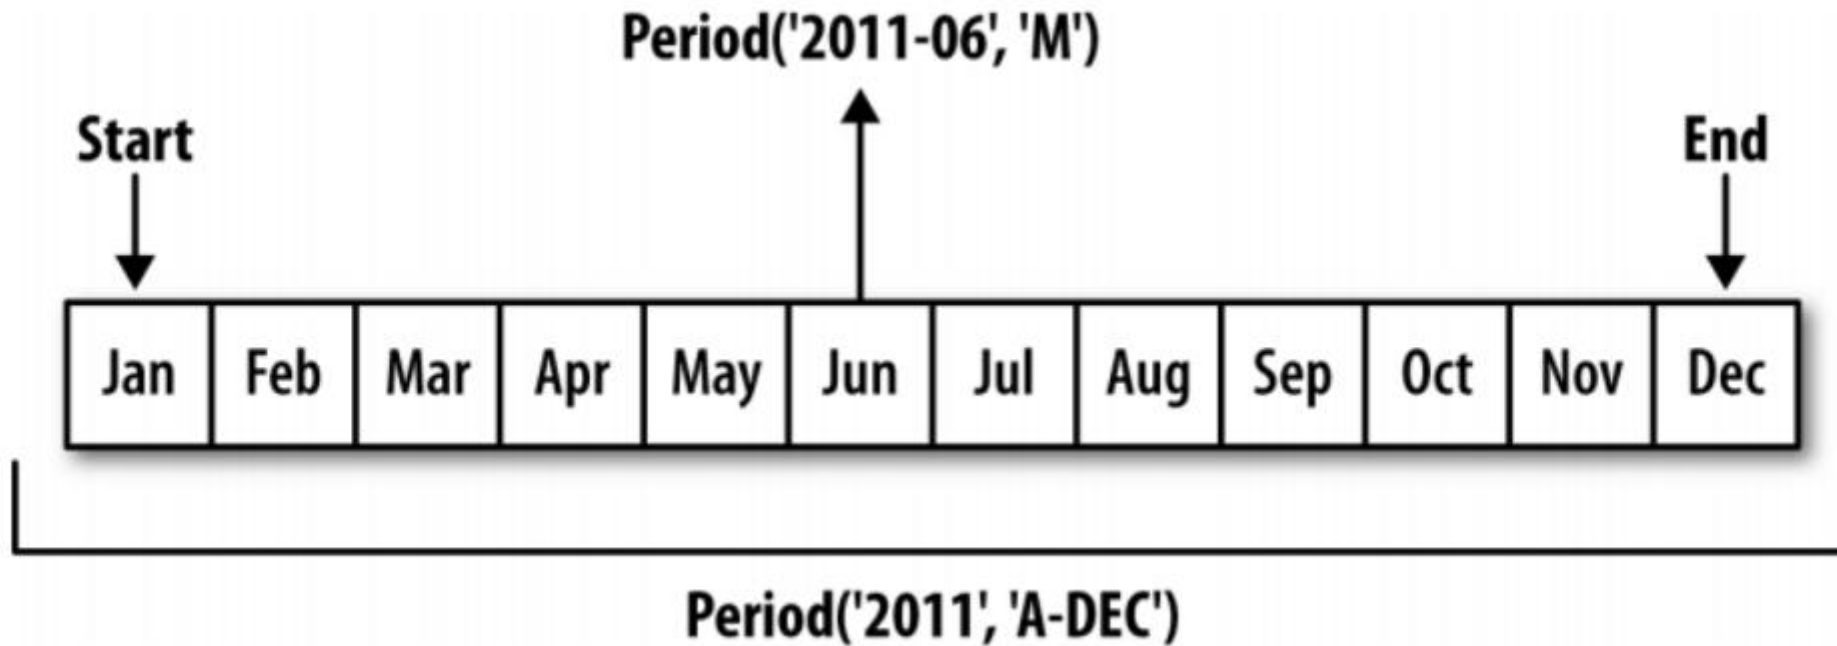
```
1 p = pd.Period('2007', freq='A-JUN')  
2 p  
3 p.asfreq('M', 'start')
```

Period('2006-07', 'M')

```
1 p.asfreq('M', 'end')
```

Period('2007-06', 'M')

Period Frequency Conversion



Period Frequency Conversion

```
1 p = pd.Period('2007-8', 'M')  
2 p.asfreq('A-JUN')
```

```
Period('2008', 'A-JUN')
```


Period Frequency Conversion

```
1 rng = pd.period_range('2006', '2009', freq='A-DEC')
2 ts = pd.Series(np.random.randn(len(rng)), index=rng)
3 ts
```

2006 -0.272657

2007 -1.692615

2008 1.423830

2009 -0.407890

Freq: A-DEC, dtype: float64

```
1 ts.asfreq('M', how='start')
```

2006-01 -0.272657

2007-01 -1.692615

2008-01 1.423830

2009-01 -0.407890

Freq: M, dtype: float64

Period Frequency Conversion

```
1 ts.asfreq('B', how='end')
```

2006-12-29 -0.272657

2007-12-31 -1.692615

2008-12-31 1.423830

2009-12-31 -0.407890

Freq: B, dtype: float64

Quarterly Period Frequencies

```
1 p = pd.Period('2012Q4', freq='Q-JAN')
2 p
```

Period('2012Q4', 'Q-JAN')

Year 2012												
M	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	DEC
Q-DEC	2012Q1			2012Q2			2012Q3			2012Q4		
Q-SEP	2012Q2			2012Q3			2012Q4			2013Q1		
Q-FEB	2012Q4		2013Q1			2013Q2			2013Q3			Q4

Different quarterly frequency conventions

Quarterly Period Frequencies

```
1 p.asfreq('D', 'start')
```

```
Period('2011-11-01', 'D')
```

```
1 p.asfreq('D', 'end')
```

```
Period('2012-01-31', 'D')
```

Quarterly Period Frequencies

```
1 p4pm = (p.asfreq('B', 'e') - 1).asfreq('T', 's') + 16 * 60
2 p4pm
```

```
Period('2012-01-30 16:00', 'T')
```

```
1 p4pm.to_timestamp()
```

```
Timestamp('2012-01-30 16:00:00')
```

Quarterly Period Frequencies

```
1 rng = pd.period_range('2011Q3', '2012Q4', freq='Q-JAN')
2 ts = pd.Series(np.arange(len(rng)), index=rng)
3 ts
```

2011Q3 0

2011Q4 1

2012Q1 2

2012Q2 3

2012Q3 4

2012Q4 5

Freq: Q-JAN, dtype: int32

Quarterly Period Frequencies

```
1 new_rng = (rng.asfreq('B', 'e') - 1).asfreq('T', 's') + 16 * 60
2 ts.index = new_rng.to_timestamp()
3 ts
```

```
2010-10-28 16:00:00    0
2011-01-28 16:00:00    1
2011-04-28 16:00:00    2
2011-07-28 16:00:00    3
2011-10-28 16:00:00    4
2012-01-30 16:00:00    5
dtype: int32
```

Converting Timestamps to Periods (and Back)

```
1 rng = pd.date_range('2000-01-01', periods=3, freq='M')
2 ts = pd.Series(np.random.randn(3), index=rng)
3 ts
```

```
2000-01-31    0.756332
2000-02-29   -1.288602
2000-03-31    0.867534
Freq: M, dtype: float64
```

```
1 pts = ts.to_period()
2 pts
```

```
2000-01    0.756332
2000-02   -1.288602
2000-03    0.867534
Freq: M, dtype: float64
```


Converting Timestamps to Periods (and Back)

```
1 rng = pd.date_range('2000/1/29', periods=6, freq='D')
2 ts2 = pd.Series(np.random.randn(6), index=rng)
3 ts2
```

```
2000-01-29    -0.252765
2000-01-30    -0.894590
2000-01-31     0.955842
2000-02-01    -1.653984
2000-02-02    -0.262528
2000-02-03    -0.976094
Freq: D, dtype: float64
```

```
1 ts2.to_period('M')
```

```
2000-01    -0.252765
2000-01    -0.894590
2000-01     0.955842
2000-02    -1.653984
2000-02    -0.262528
2000-02    -0.976094
Freq: M, dtype: float64
```

Converting Timestamps to Periods (and Back)

```
1 pts = ts2.to_period()  
2 pts
```

```
2000-01-29    -0.252765  
2000-01-30    -0.894590  
2000-01-31     0.955842  
2000-02-01    -1.653984  
2000-02-02    -0.262528  
2000-02-03    -0.976094  
Freq: D, dtype: float64
```

```
1 pts.to_timestamp(how='end')
```

```
2000-01-29 23:59:59.999999999    -0.252765  
2000-01-30 23:59:59.999999999    -0.894590  
2000-01-31 23:59:59.999999999     0.955842  
2000-02-01 23:59:59.999999999    -1.653984  
2000-02-02 23:59:59.999999999    -0.262528  
2000-02-03 23:59:59.999999999    -0.976094  
Freq: D, dtype: float64
```

Creating a PeriodIndex from Arrays

```
1 data = pd.read_csv('macrodata.csv')  
2 data.head(5)
```

	year	quarter	realgdp	realcons	realinv	realgovt	realdpi	cpi	m1	tbilrate	unemp
0	1959.0	1.0	2710.349	1707.4	286.898	470.045	1886.9	28.98	139.7	2.82	
1	1959.0	2.0	2778.801	1733.7	310.859	481.301	1919.7	29.15	141.7	3.08	
2	1959.0	3.0	2775.488	1751.8	289.226	491.260	1916.4	29.35	140.5	3.82	
3	1959.0	4.0	2785.204	1753.7	299.356	484.052	1931.3	29.37	140.0	4.33	
4	1960.0	1.0	2847.699	1770.5	331.722	462.199	1955.5	29.54	139.6	3.50	

Creating a PeriodIndex from Arrays

1	data.year
---	-----------

0	1959.0
1	1959.0
2	1959.0
3	1959.0
4	1960.0
5	1960.0

...

200	2009.0
201	2009.0
202	2009.0

Name: year, Length: 203, dtype: float64

1	data.quarter
---	--------------

0	1.0
1	2.0
2	3.0
3	4.0
4	1.0
5	2.0

...

200	1.0
201	2.0
202	3.0

Name: quarter, Length: 203, dtype: float64

Creating a PeriodIndex from Arrays

```
1 index = pd.PeriodIndex(year=data.year, quarter=data.quarter,  
2                        freq='Q-DEC')  
3 index
```

```
PeriodIndex(['1959Q1', '1959Q2', '1959Q3', '1959Q4', '1960Q1', '1960Q2',  
            '1960Q3', '1960Q4', '1961Q1', '1961Q2',  
            ...,  
            '2007Q2', '2007Q3', '2007Q4', '2008Q1', '2008Q2', '2008Q3',  
            '2008Q4', '2009Q1', '2009Q2', '2009Q3'],  
            dtype='period[Q-DEC]', length=203, freq='Q-DEC')
```

Creating a PeriodIndex from Arrays

```
1 data.index = index
2 data.infl
```

1959Q1	0.00
1959Q2	2.34
1959Q3	2.74
1959Q4	0.27
1960Q1	2.31
1960Q2	0.14
...	
2008Q4	-8.79
2009Q1	0.94
2009Q2	3.37
2009Q3	3.56

Freq: Q-DEC, Name: infl, Length: 203, dtype: float64

Resampling and Frequency Conversion

```
1 rng = pd.date_range('2000-01-01', periods=100, freq='D')
2 ts = pd.Series(np.random.randn(len(rng)), index=rng)
3 ts
```

2000-01-01 -1.493407

2000-01-02 1.167858

2000-01-03 0.969001

2000-01-04 -2.536487

2000-01-05 0.362754

...

2000-04-08 -1.040816

2000-04-09 0.426419

Freq: D, Length: 100, dtype: float64

Resampling and Frequency Conversion

```
1 ts.resample('M').mean()
```

```
2000-01-31    0.027804
2000-02-29    0.194619
2000-03-31    0.166734
2000-04-30    0.239600
Freq: M, dtype: float64
```

```
1 ts.resample('M', kind='period').mean()
```

```
2000-01    0.027804
2000-02    0.194619
2000-03    0.166734
2000-04    0.239600
Freq: M, dtype: float64
```


Resampling and Frequency Conversion

resample method arguments

Argument	Description
<code>freq</code>	String or DateOffset indicating desired resampled frequency, e.g. 'M', '5min', or <code>Second(15)</code>
<code>how='mean'</code>	Function name or array function producing aggregated value, for example 'mean', 'ohlc', <code>np.max</code> . Defaults to 'mean'. Other common values: 'first', 'last', 'median', 'ohlc', 'max', 'min'.
<code>axis=0</code>	Axis to resample on, default <code>axis=0</code>
<code>fill_method=None</code>	How to interpolate when upsampling, as in 'ffill' or 'bfill'. By default does no interpolation.
<code>closed='right'</code>	In downsampling, which end of each interval is closed (inclusive), 'right' or 'left'. Defaults to 'right'
<code>label='right'</code>	In downsampling, how to label the aggregated result, with the 'right' or 'left' bin edge. For example, the 9:30 to 9:35 5-minute interval could be labeled 9:30 or 9:35. Defaults to 'right' (or 9:35, in this example).

Resampling and Frequency Conversion

resample method arguments

Argument	Description
<code>loffset=None</code>	Time adjustment to the bin labels, such as <code>'-1s' / Second(-1)</code> to shift the aggregate labels one second earlier
<code>limit=None</code>	When forward or backward filling, the maximum number of periods to fill
<code>kind=None</code>	Aggregate to periods (<code>'period'</code>) or timestamps (<code>'timestamp'</code>); defaults to kind of index the time series has
<code>convention=None</code>	When resampling periods, the convention (<code>'start'</code> or <code>'end'</code>) for converting the low frequency period to high frequency. Defaults to <code>'end'</code>

Downsampling

```
1 rng = pd.date_range('2000-01-01', periods=12, freq='T')
2 ts = pd.Series(np.arange(12), index=rng)
3 ts
```

```
2000-01-01 00:00:00    0
2000-01-01 00:01:00    1
2000-01-01 00:02:00    2
2000-01-01 00:03:00    3
...
2000-01-01 00:09:00    9
2000-01-01 00:10:00   10
2000-01-01 00:11:00   11
Freq: T, dtype: int32
```

Downsampling

- Which side of each interval is *closed*
- How to label each aggregated bin, either with the start of the interval or the end

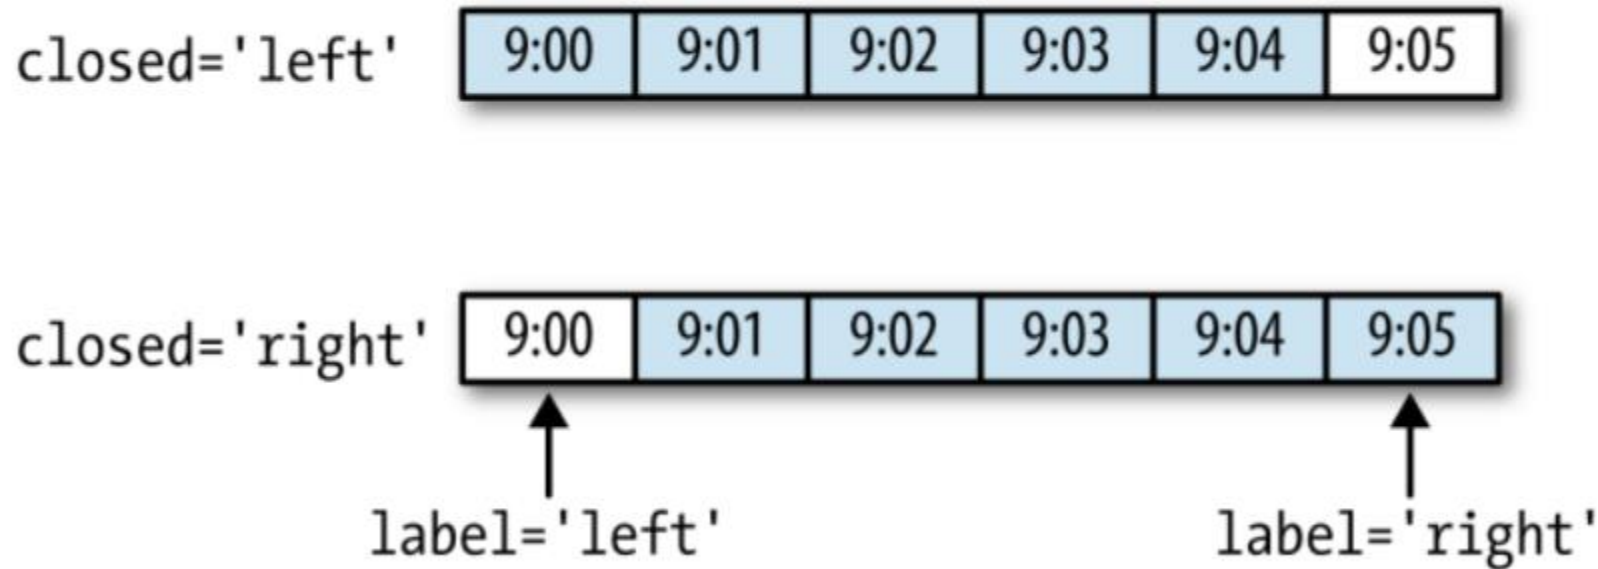
```
1 ts.resample('5min', closed='right')
```

DatetimeIndexResampler [freq=<5 * Minutes>, axis=0, closed=right, label=left, convention=start, base=0]

```
1 ts.resample('5min', closed='right').sum()
```

```
1999-12-31 23:55:00    0
2000-01-01 00:00:00   15
2000-01-01 00:05:00   40
2000-01-01 00:10:00   11
Freq: 5T, dtype: int32
```

Downsampling



minute resampling of closed, label conventions

Downsampling

```
1 ts.resample('5min', closed='right', label='right').sum()
```

```
2000-01-01 00:00:00    0
2000-01-01 00:05:00   15
2000-01-01 00:10:00   40
2000-01-01 00:15:00   11
Freq: 5T, dtype: int32
```

```
1 ts.resample('5min', closed='right',
2               label='right', loffset='-1s').sum()
```

```
1999-12-31 23:59:59    0
2000-01-01 00:04:59   15
2000-01-01 00:09:59   40
2000-01-01 00:14:59   11
Freq: 5T, dtype: int32
```

Open-High-Low-Close (OHLC) resampling

```
1 ts.resample('5min').ohlc()
```

	open	high	low	close
2000-01-01 00:00:00	0	4	0	4
2000-01-01 00:05:00	5	9	5	9
2000-01-01 00:10:00	10	11	10	11

Upsampling and Interpolation

```
1 frame = pd.DataFrame(np.random.randn(2, 4),  
2                       index=pd.date_range('2000/1/1', periods=2,  
3                                           freq='W-WED'),  
4                       columns=['Colorado', 'Texas', 'New York', 'Ohio'])  
5 frame
```

	Colorado	Texas	New York	Ohio
2000-01-05	2.268799	0.146326	0.508391	-0.196713
2000-01-12	-3.745356	-1.520113	-0.346839	-0.696918

Upsampling and Interpolation

```
1 df_daily = frame.resample('D').asfreq()  
2 df_daily
```

	Colorado	Texas	New York	Ohio
2000-01-05	2.268799	0.146326	0.508391	-0.196713
2000-01-06	NaN	NaN	NaN	NaN
2000-01-07	NaN	NaN	NaN	NaN
2000-01-08	NaN	NaN	NaN	NaN
2000-01-09	NaN	NaN	NaN	NaN
2000-01-10	NaN	NaN	NaN	NaN
2000-01-11	NaN	NaN	NaN	NaN
2000-01-12	-3.745356	-1.520113	-0.346839	-0.696918

Upsampling and Interpolation

```
1 frame.resample('D').ffill()
```

	Colorado	Texas	New York	Ohio
2000-01-05	2.268799	0.146326	0.508391	-0.196713
2000-01-06	2.268799	0.146326	0.508391	-0.196713
2000-01-07	2.268799	0.146326	0.508391	-0.196713
2000-01-08	2.268799	0.146326	0.508391	-0.196713
2000-01-09	2.268799	0.146326	0.508391	-0.196713
2000-01-10	2.268799	0.146326	0.508391	-0.196713
2000-01-11	2.268799	0.146326	0.508391	-0.196713
2000-01-12	-3.745356	-1.520113	-0.346839	-0.696918

Upsampling and Interpolation

```
1 frame.resample('D').ffill(limit=2)
```

	Colorado	Texas	New York	Ohio
2000-01-05	2.268799	0.146326	0.508391	-0.196713
2000-01-06	2.268799	0.146326	0.508391	-0.196713
2000-01-07	2.268799	0.146326	0.508391	-0.196713
2000-01-08	NaN	NaN	NaN	NaN
2000-01-09	NaN	NaN	NaN	NaN
2000-01-10	NaN	NaN	NaN	NaN
2000-01-11	NaN	NaN	NaN	NaN
2000-01-12	-3.745356	-1.520113	-0.346839	-0.696918

Upsampling and Interpolation

```
1 frame.resample('W-THU')
```

DatetimeIndexResampler [freq=<Week: weekday=3>, axis=0, closed=right, label=right, convention=start, base=0]

```
1 frame.resample('W-THU').ffill()
```

	Colorado	Texas	New York	Ohio
2000-01-06	2.268799	0.146326	0.508391	-0.196713
2000-01-13	-3.745356	-1.520113	-0.346839	-0.696918

Resampling with Periods

```
1 frame = pd.DataFrame(np.random.randn(24, 4),
2                       index=pd.period_range('2000-1', '2001-12',
3                                             freq='M'),
4                       columns=['Colorado', 'Texas', 'New York', 'Ohio'])
5 frame[:5]
```

	Colorado	Texas	New York	Ohio
2000-01	0.873921	-1.180212	-0.208885	-0.549671
2000-02	-1.252880	-1.276761	1.881156	1.108227
2000-03	-1.751994	-0.973899	0.908732	-0.509226
2000-04	-1.023400	-0.412273	-1.073039	-0.601411
2000-05	0.222178	0.949363	0.704186	-1.358964

Resampling with Periods

```
1 annual_frame = frame.resample('A-DEC').mean()  
2 annual_frame
```

	Colorado	Texas	New York	Ohio
2000	0.154403	-0.402994	-0.115005	0.032294
2001	0.326664	0.377224	-0.324770	-0.607382

Resampling with Periods

```
1 # Q-DEC: Quarterly, year ending in December
2 annual_frame.resample('Q-DEC')
```

PeriodIndexResampler [freq=<QuarterEnd: startingMonth=12>, axis=0, closed=right, label=right, convention=start, base=0]

```
1 annual_frame.resample('Q-DEC').ffill()
```

	Colorado	Texas	New York	Ohio
2000Q1	0.154403	-0.402994	-0.115005	0.032294
2000Q2	0.154403	-0.402994	-0.115005	0.032294
2000Q3	0.154403	-0.402994	-0.115005	0.032294
2000Q4	0.154403	-0.402994	-0.115005	0.032294

2001Q1	0.326664	0.377224	-0.324770	-0.607382
2001Q2	0.326664	0.377224	-0.324770	-0.607382
2001Q3	0.326664	0.377224	-0.324770	-0.607382
2001Q4	0.326664	0.377224	-0.324770	-0.607382

Resampling with Periods

```
1 annual_frame.resample('Q-DEC', convention='end').ffill()
```

	Colorado	Texas	New York	Ohio
2000Q4	0.154403	-0.402994	-0.115005	0.032294
2001Q1	0.154403	-0.402994	-0.115005	0.032294
2001Q2	0.154403	-0.402994	-0.115005	0.032294
2001Q3	0.154403	-0.402994	-0.115005	0.032294
2001Q4	0.326664	0.377224	-0.324770	-0.607382

Resampling with Periods

- In downsampling, the target frequency must be a *subperiod* of the source frequency.
- In upsampling, the target frequency must be a *superperiod* of the source frequency.

```
1 annual_frame.resample('Q-MAR').ffill()
```

	Colorado	Texas	New York	Ohio
2000Q4	0.154403	-0.402994	-0.115005	0.032294
2001Q1	0.154403	-0.402994	-0.115005	0.032294
2001Q2	0.154403	-0.402994	-0.115005	0.032294
2001Q3	0.154403	-0.402994	-0.115005	0.032294
2001Q4	0.326664	0.377224	-0.324770	-0.607382

2002Q1	0.326664	0.377224	-0.324770	-0.607382
2002Q2	0.326664	0.377224	-0.324770	-0.607382
2002Q3	0.326664	0.377224	-0.324770	-0.607382